

# Flight Price Analytics — Project Report

## 1) Project Overview

Built an end-to-end analytics stack to answer a traveler's core question: "**Buy now or wait?**" The solution ingests daily fare snapshots, cleans and models them, and exposes **actionable pricing signals** (lead-time sweet spots, weekday vs. weekend effects, and route alerts) through a **one-page Power BI dashboard**. Structure and deliverable style mirror the author's prior retail analytics project to keep outputs recruiter-friendly and decision-ready.

## 2) Dataset Summary

- **Grain:** one row per (route, search\_date, depart\_date) fare snapshot
- **Span:** ~Jun-2025 → Nov-2025 (sample shown; scalable to longer history)
- **Rows / Cols:** 133776 rows and 9 columns
- **Core fields:** origin, destination, route, price, search\_date, depart\_date, lead\_time\_days
- **Derived fields:** route key, **lead-time bucket** (0–7, 8–14, 15–30, 31–60, 60+), **is\_weekend** (by depart\_date), 7-day baseline of price for forecast vs. actual, and 7-day % change for alerting

## 3) Exploratory Data Analysis (Python)

**Stack:** pandas / numpy / matplotlib.

### Steps

- **Health:** .info(), NA scan, duplicate scan at (route, depart\_date, snapshot\_date); clipped non-positive prices.

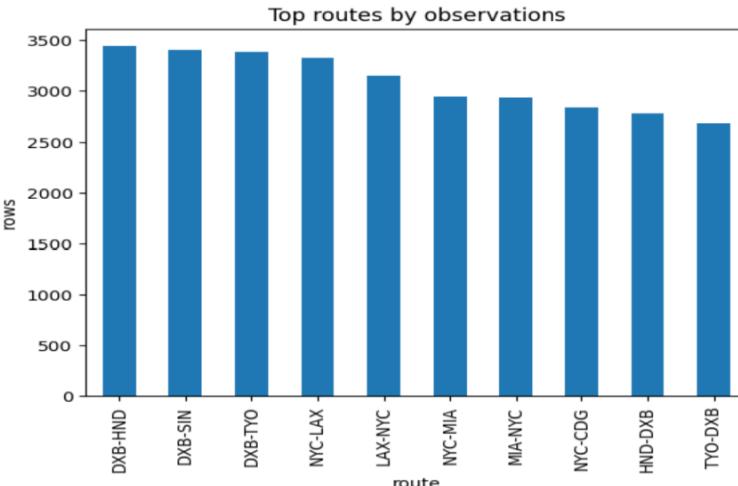
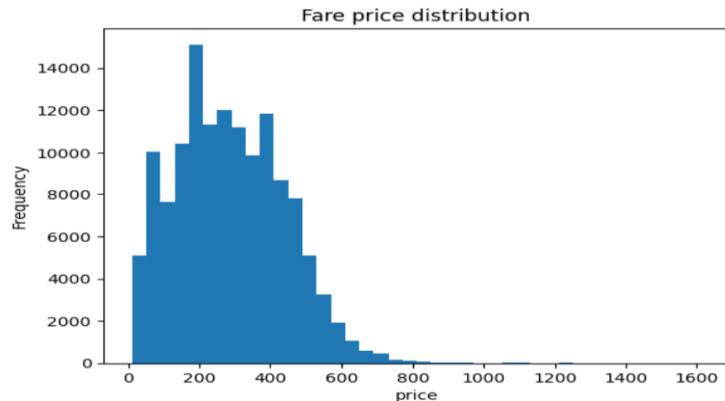
```
[4]: df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 133775 entries, 0 to 133774
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   source_name  133775 non-null  object  
 1   origin       133775 non-null  object  
 2   destination  133775 non-null  object  
 3   search_date  133775 non-null  object  
 4   depart_date  133775 non-null  object  
 5   price        133775 non-null  int64  
dtypes: int64(1), object(5)
memory usage: 6.1+ MB
```

- **Feature engineering:**

	source_name	origin	destination	search_date	depart_date	price	snapshot_date	route	lead_time_days
0	flight_prices_2025-06-07.csv	NYC	LHR	2025-06-07	2025-08-01	140	2025-06-07	NYC-LHR	55
1	flight_prices_2025-06-07.csv	NYC	LHR	2025-06-07	2025-08-25	169	2025-06-07	NYC-LHR	79
2	flight_prices_2025-06-07.csv	NYC	LHR	2025-06-07	2025-08-04	172	2025-06-07	NYC-LHR	58
3	flight_prices_2025-06-07.csv	NYC	LHR	2025-06-07	2025-07-28	180	2025-06-07	NYC-LHR	51
4	flight_prices_2025-06-07.csv	NYC	LHR	2025-06-07	2025-07-21	184	2025-06-07	NYC-LHR	44

- $\text{lead\_time\_days} = (\text{depart\_date} - \text{search\_date})$
- lead\_bucket via SWITCH/CASE bands for charting
- $\text{is\_weekend} = \text{WEEKDAY}(\text{depart\_date}, 2) \geq 6$  (Sat/Sun)
- 7-day rolling **baseline** of price to compare against last snapshot
- Quick EDA (by route):

```
# Price distribution
df["price"].plot(kind="hist", bins=40, title="Fare price distribution")
plt.xlabel("price"); plt.show()
```



- **Seasonality:** visible month-to-month waves in average search-day price
- **Lead-time curve:** some routes show **U-shaped** price vs. lead time (cheap around 15–30 days, higher very early/very late)
- **Weekend effect:** for several routes, **weekend mean < weekday mean.**

## 4) Data Analysis using SQL

Reusable SQL answered the business questions below (snippets are in the project SQL file).

--1) Which routes are the most price-volatile?

--Why: volatile routes are good candidates for alerts and “Buy now” nudges.

	route text	days bigint	avg_price numeric	stdev numeric
1	HND-MIA	149	631.88	150.72
2	TYO-MIA	149	551.44	141.79
3	MIA-SYD	103	700.89	137.09
4	SYD-CDG	155	544.56	92.54
5	SYD-MIA	85	597.74	91.63
6	MIA-SIN	75	502.48	89.31
7	MIA-HND	154	570.23	77.55
8	NYC-SYD	155	562.46	68.40
9	MIA-TYO	154	476.40	61.84
10	HND-NYC	155	492.58	61.60

--2) What does the lead-time price curve look like for a route?

	lead_band text	n bigint	avg_price numeric
1	00–07	407	157.13
2	08–14	409	156.09
3	15–21	413	156.00
4	22–28	371	154.27
5	29–42	747	154.29
6	43–60	602	155.94
7	61+	412	156.02

--3) “Buy or Wait” for a given route & departure?

--Compare **latest snapshot**'s price to the last 30 days of snapshots (same route+depart\_date). Flag if latest is  $\geq 10\%$  below the 30-day median.

A screenshot of a database interface showing a single row of data in a table. The table has three columns: 'latest\_price' (datatype bigint), 'median\_30d' (datatype numeric), and 'recommendation' (datatype text). The data row shows values: 153, 154.00, and 'WAIT' respectively. The 'recommendation' column has a lock icon.

	latest_price bigint	median_30d numeric	recommendation text
1	153	154.00	WAIT

--4) What day of week is typically cheapest to depart (by route)?

A screenshot of a database interface showing a list of routes and their average prices by departure day of the week. The table has five columns: 'route' (datatype text), 'depart\_dow' (datatype text), 'avg\_price' (datatype numeric), and 'n' (datatype bigint). The data shows various routes and their average prices for different days of the week, along with the count of flights. The 'route' column has a lock icon.

	route text	depart_dow text	avg_price numeric	n bigint
1	CDG-DXB	Fri	192.14	404
2	CDG-DXB	Mon	202.28	393
3	CDG-DXB	Sat	198.72	249
4	CDG-DXB	Sun	197.26	383
5	CDG-DXB	Thu	197.53	113
6	CDG-DXB	Tue	196.54	119
7	CDG-DXB	Wed	188.22	374
8	CDG-HND	Fri	447.37	201
9	CDG-HND	Mon	447.27	345
10	CDG-HND	Sat	448.76	139
11	CDG-HND	Sun	452.38	110
12	CDG-HND	Thu	449.46	356
13	CDG-HND	Tue	447.16	312

Total rows: 614    Query complete 00:00:00.159

--5) What booking window (lead time) tends to have the lowest average price per route?

The screenshot shows a database query interface with a toolbar at the top containing icons for new table, file, copy, paste, delete, download, and SQL. Below the toolbar is a table with the following data:

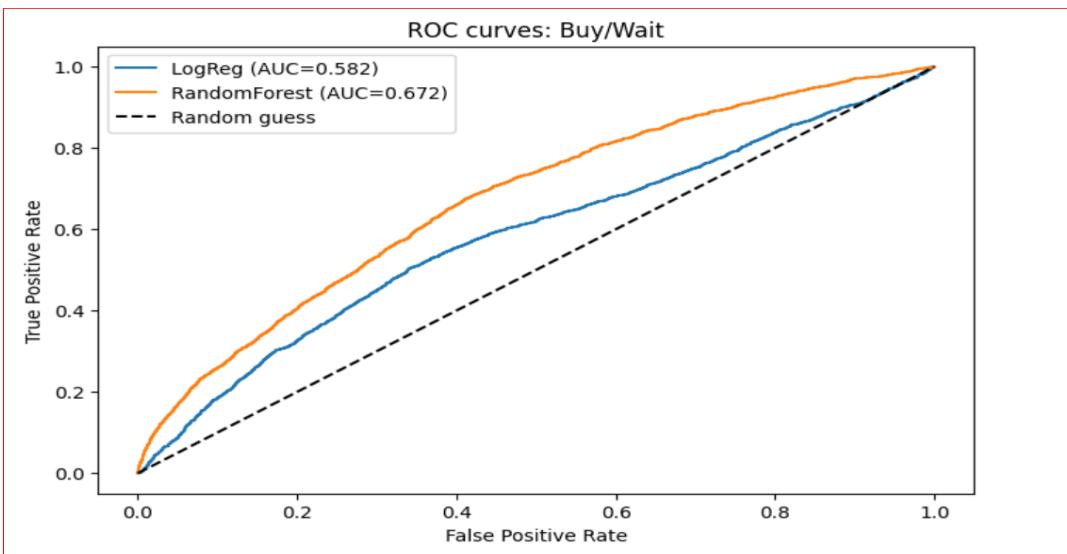
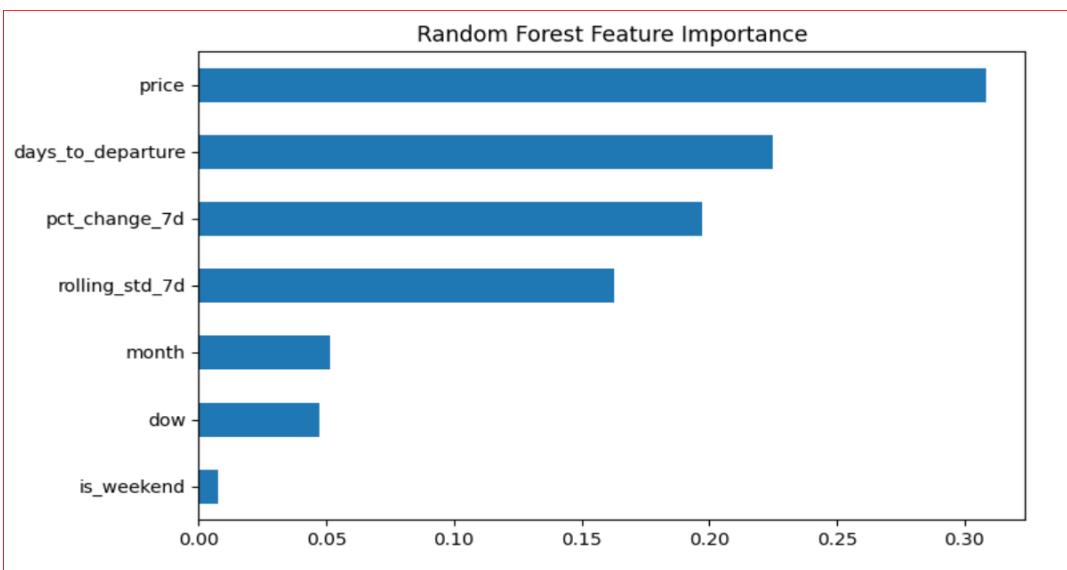
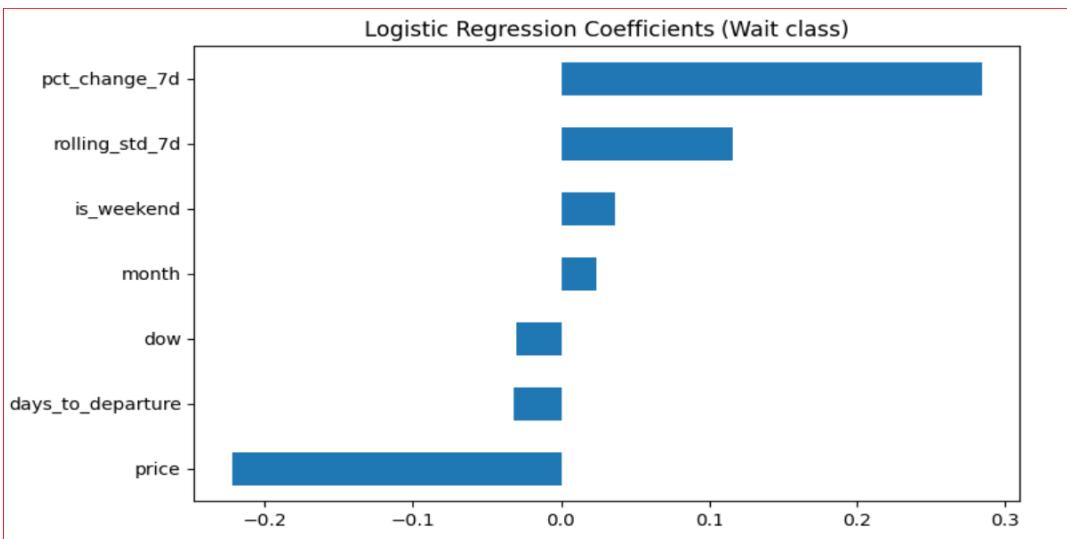
	route text	best_lead_time_days	best_avg_price	n
1	CDG-DXB	40	181.38	21
2	CDG-HND	22	426.25	16
3	CDG-LAX	66	304.74	19
4	CDG-LHR	10	80.00	1
5	CDG-MIA	73	325.20	10
6	CDG-NYC	84	266.23	22
7	CDG-SIN	100	325.00	1
8	CDG-SYD	52	558.54	13
9	CDG-TYO	22	424.45	20
10	DXB-CDG	38	176.72	43
11	DXB-HND	70	292.72	36
12	DXB-LAX	98	467.90	10
13	DYBLHD	88	160.75	1

Total rows: 88    Query complete 00:00:00.158

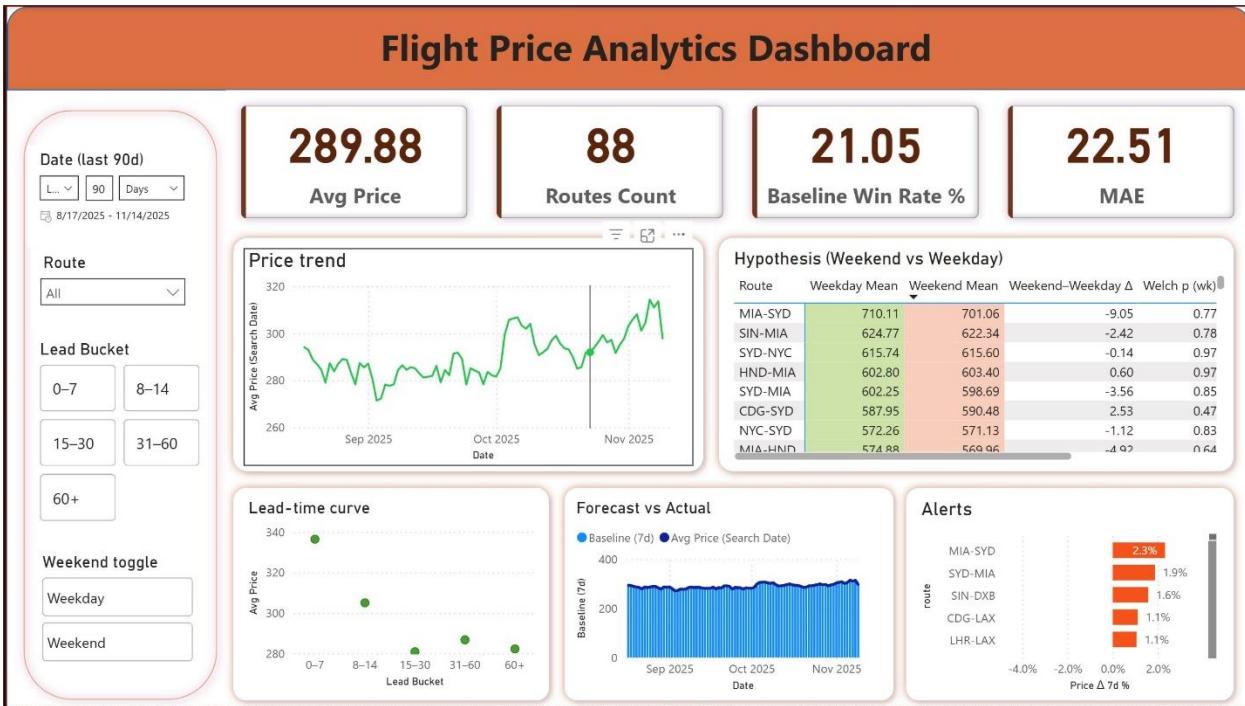
This modular, question-first query pattern follows the same “analysis menu” approach used in the floral project to support BI and ad-hoc analysis.

## 5) Buy/Wait ML Model (Python, scikit-learn)

Framed “**Buy now vs Wait**” as a **supervised classification** problem using engineered features (current price vs rolling **7-day mean**, volatility, days-to-departure, DOW/month, weekend flag). Trained **logistic regression** and class-balanced **random forest** models with time-based **cross-validation**; evaluated with **ROC AUC**, precision/recall, and confusion matrices. Exposed a simple `buy_or_wait()` helper that returns both the decision and the probability of “Wait” for use in the app and alerts.



## 5) Dashboard in Power BI



## 6) Business Recommendations

### A. Daily “Buy” list

- Compare today's price to the **30-day median**.
- If today is **≥10% cheaper**, flag the route/departure as **BUY** and surface it in email/banner.

### B. When to book (per route)

- Publish each route's **best lead-time window** (e.g., “15–30 days out”).
- Refresh monthly so guidance stays current.

### C. Weekend messaging

- If **weekend < weekday** prices, promote “**Save more on weekend departures!**”
- Time ads and push notifications on **Thu–Fri**.

### D. Volatility watchlist

- Track routes with the **highest price variance**.

- Add **price-drop badges** in search results and poll those routes more frequently.

#### E. Keep improving the signals

- Monitor **Baseline Win Rate** and **MAE**; switch to longer, seasonal baselines (e.g., **28-day**) where errors stay high.
- Retain the **question-driven SQL menu + one-page dashboard** so teams can self-serve and act fast.

#### F. ML-powered Buy/Wait guidance

- Use the random-forest Buy/Wait model alongside simple rules to score each route/departure.
- Surface high-confidence “Wait” and “Buy now” recommendations in the dashboard and app, and track model performance (ROC AUC, precision/recall) over time.